

The Big Picture

Index

1. Objectives
2. Deep Learning
3. Understanding TensorFlow
4. 'Hello World!'



Index

1. **Objectives**
2. Deep Learning
3. Understanding TensorFlow
4. 'Hello World!'

Objectives

- We are going to address the next questions:
 - What kind of problems are we going to deal with?
 - Why do we need tools such a TF?
 - How to use TF? First steps...




Index

1. Objectives
2. **Deep Learning**
3. Understanding TensorFlow
4. 'Hello World!'

Deep Learning

- Deep Learning algorithms can be described as an instance of a simple recipe:
 - Dataset (MNIST, CIFAR, BRATS...)
 - Model (ANN, ConvNN, RNN...)
 - Cost Function (MSE, Cross Entropy, DICE...)
 - Optimization Procedure (1st order GD/LM, 2nd order Newton...)

Deep Learning




X^m

$$F : (X^m, W) \mapsto Y_p^n$$

$$J : (Y_t^n, Y_p^n) \mapsto \mathbb{R}$$

$$\frac{\partial J}{\partial W}$$





Deep Learning

$$X^m$$

i) DEL ORDEN DE
MILLONES!
(CON SUERTE)

$$F : (X^m, W) \mapsto Y_p^n$$


ii) DEL ORDEN DE
CIENTOS DE
MILLONES!!

$$J : (Y_t^n, Y_p^n) \mapsto \mathbb{R}$$

$$\frac{\partial J}{\partial W}$$

iii) LA DERIVADA PARA
CADA UNO DE LOS
PESOS!!!

iiii) Y ENCIMA ES UN PROCESO ITERATIVO!!!!



Index

1. Objectives
2. Deep Learning
- 3. Understanding TensorFlow**
4. 'Hello World!'

Understanding TensorFlow

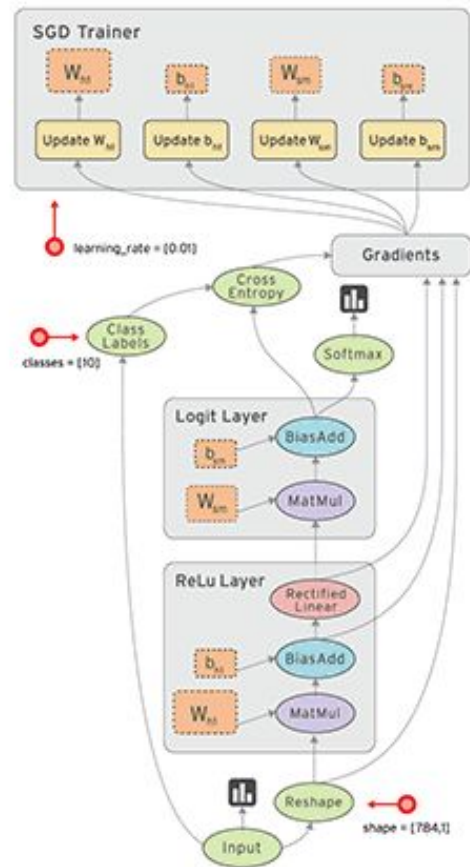
TensorFlow allows us to handle easily the computational cost and the calculations of the derivatives.

- TF works with Tensors, and all the tensors build a **computational graph**.
- TF uses **automatic differentiation** in order to calculate the derivatives of each parameter. Therefore, each node of the graph knows how to compute its own gradient.

Understanding TensorFlow

We might think TF programs as consisting of two sections:

1. Building the computational graph
2. Running the graph.



Understanding TensorFlow

1_example.py




The Tensor primitive (Extension of *Eigen* Tensor class)

```
1  3 # a rank 0 tensor; this is a scalar with shape []
2  [1., 2., 3.] # a rank 1 tensor; this is a vector with shape [3]
3  [[1., 2., 3.], [4., 5., 6.]] # a rank 2 tensor; a matrix with shape [2, 3]
4  [[[1., 2., 3.]], [[7., 8., 9.]]] # a rank 3 tensor with shape [2, 1, 3]
```

Building the computational graph: Tensors and Operations

```
1  import tensorflow as tf
2
3  node1 = tf.constant(3.0, dtype=tf.float32, name='node1')
4  node2 = tf.constant(4.0, dtype=tf.float32, name='node2')
5  node3 = tf.add(node1, node2)
6
7  print(node1, node3)
8
```



```
(ibime) ger@devon:~/Desktop/TF$ python example1.py
Tensor("node1:0", shape=(), dtype=float32) Tensor("Add:0", shape=(), dtype=float32)
```

Understanding TensorFlow

1_example.py



Running the computational graph

```
1 import tensorflow as tf
2
3 node1 = tf.constant(3.0, dtype=tf.float32, name='node1')
4 node2 = tf.constant(4.0, dtype=tf.float32, name='node2')
5 node3 = tf.add(node1, node2)
6
7 print(node1, node3)
8
```


```
9 sess = tf.Session()
10 result = sess.run([node3])
11 print('Result:', result)
```

```
Result: [7.0]
(ibime) ger@devon:~/Desktop/TF$
```



QUIERO ESO

Understanding TensorFlow



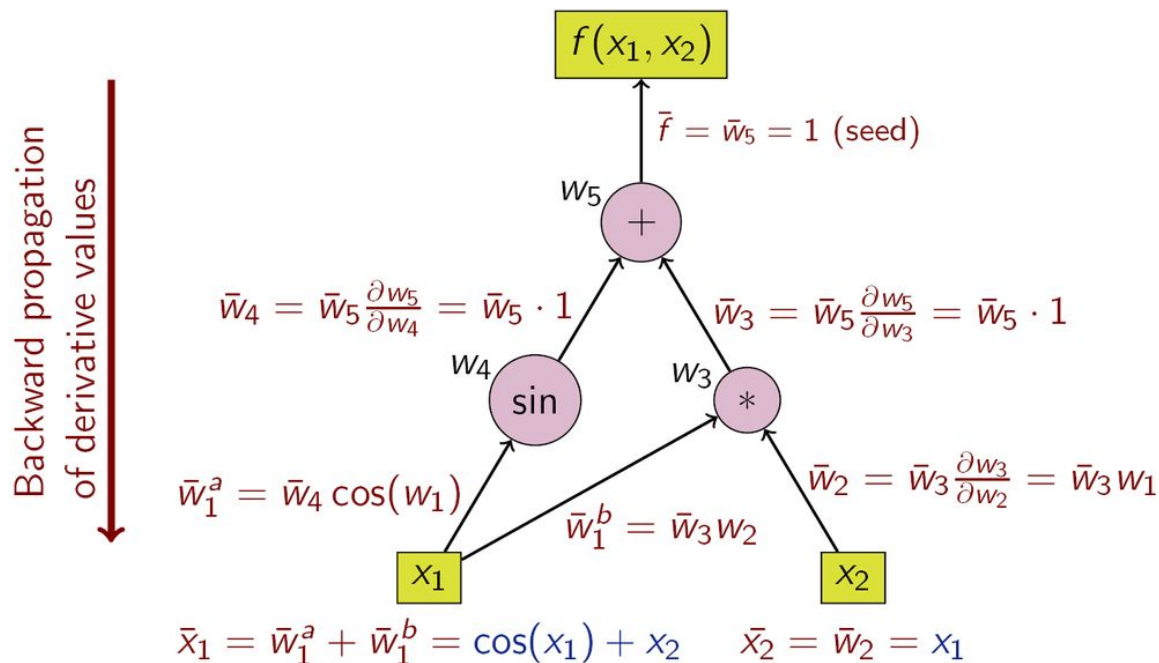
1) LOAD THE DATA

2) DEFINE THE MODEL BUILDING THE COMPUTATIONAL GRAPH

3) TRAIN THE MODEL RUNNING THE GRAPH (SESSION)

Understanding TensorFlow

What does it mean **automatic differentiation**?



Contents

Define the op's interface

Implement the kernel for the op

Multi-threaded CPU kernels

GPU kernels

Build the op library

Compile the op using your system compiler (TensorFlow binary installation)

Compile the op using bazel (TensorFlow source installation)

Use the op in Python

Verify that the op works

Building advanced features into your op

Conditional checks and validation

Op registration

GPU Support

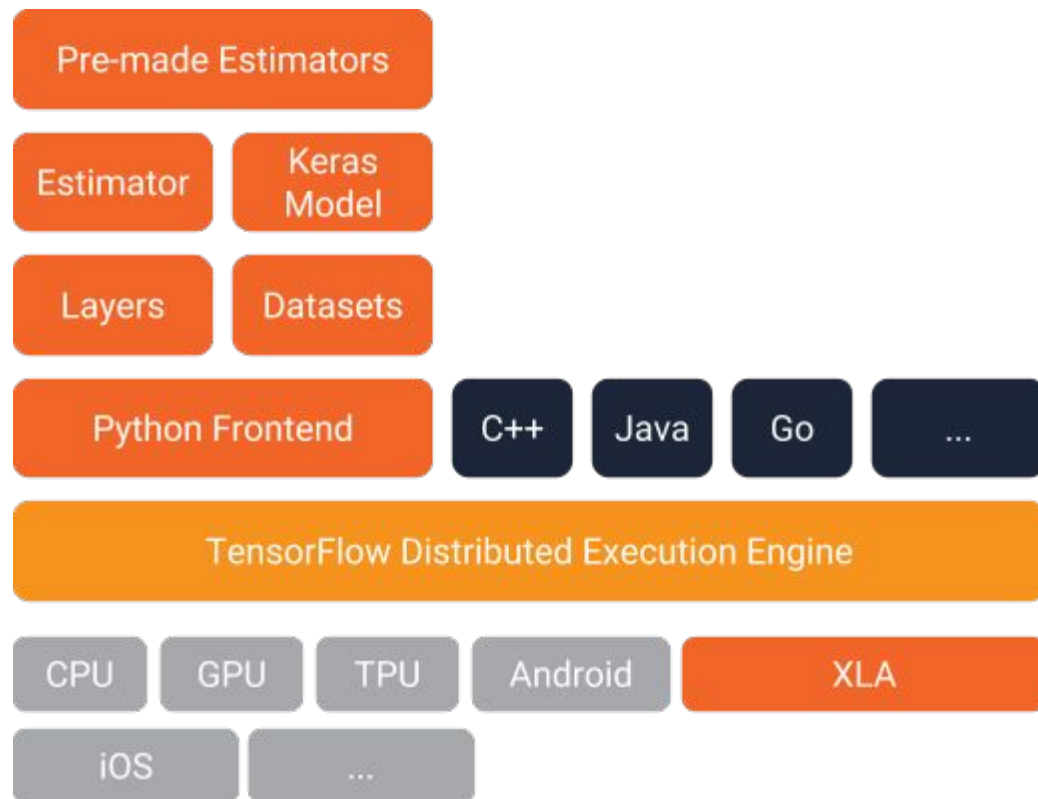
Implement the gradient in Python

Shape functions in C++

Understanding TensorFlow

Summary:

- TensorFlow is a low level library that allow us to implement a bunch of DeepLearning models efficiently.

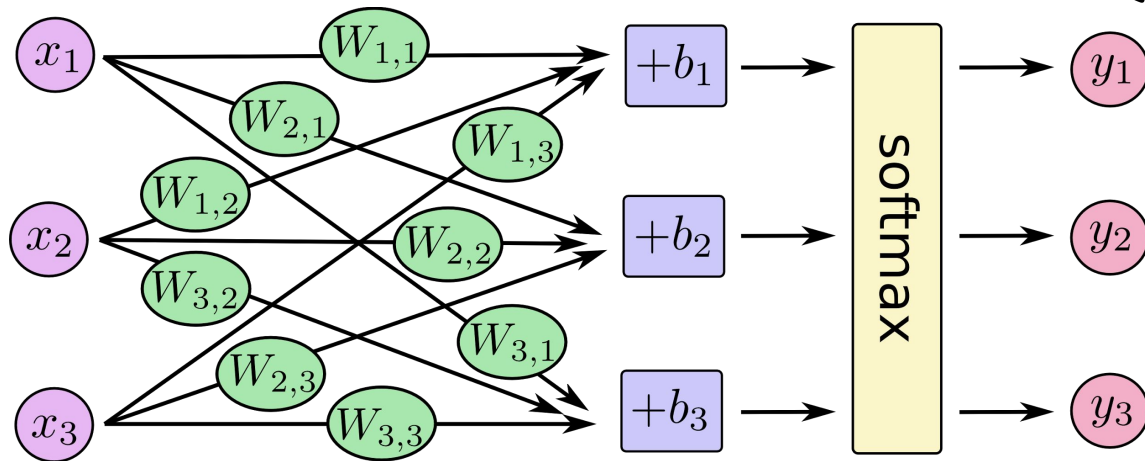



Index

1. Objectives
2. Deep Learning
3. Understanding TensorFlow
4. **'Hello World!'**

'Hello World!'

2_example.py



$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \end{bmatrix} = \text{softmax} \left(\begin{bmatrix} W_{1,1} & W_{1,2} & W_{1,3} \\ W_{2,1} & W_{2,2} & W_{2,3} \\ W_{3,1} & W_{3,2} & W_{3,3} \end{bmatrix} \cdot \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} + \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix} \right) \quad y_p = \text{softmax}(W \cdot x + b)$$

'Hello World!'

2_example.py



1 1 5 4 3
7 5 3 5 3
5 5 9 0 6
3 5 2 0 0

$$y_p = \text{softmax}(W \cdot x + b)$$

```
1 import tensorflow as tf
2 from tensorflow.examples.tutorials.mnist import input_data

8 # Load the data #
9 mnist = input_data.read_data_sets("/tmp/data/", one_hot=True)
10 x_train, y_train = mnist.train.next_batch(batch_size)
11

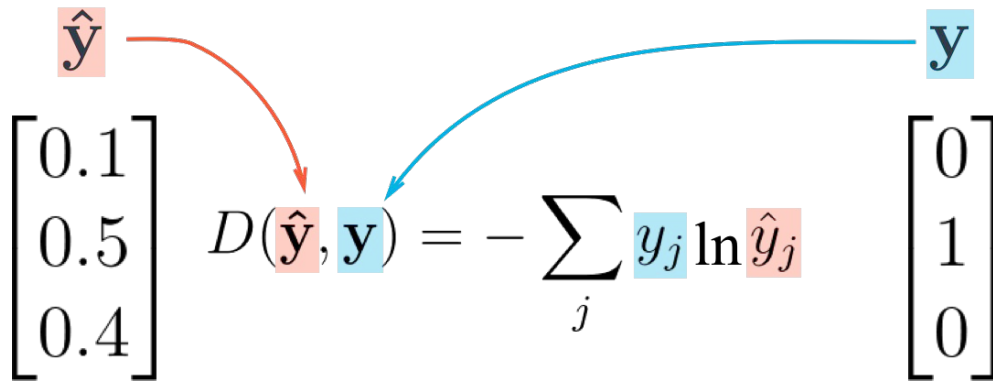
15 # Declare parameters #
16 weights = {
17     'w1': tf.Variable(tf.random_normal([784, 10]))
18 }
19 biases = {
20     'b1': tf.Variable(tf.random_normal([10]))
21 }
22
```

'Hello World!'

2_example.py



Cost Function


$$D(\hat{\mathbf{y}}, \mathbf{y}) = - \sum_j y_j \ln \hat{y}_j$$

Left vector ($\hat{\mathbf{y}}$): $\begin{bmatrix} 0.1 \\ 0.5 \\ 0.4 \end{bmatrix}$

Right vector (\mathbf{y}): $\begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$


```
31 # Cost Function. Compare predictions with true labels
32 cross_entropy = tf.reduce_mean([- y_pred * tf.log(y_in)])
```

'Hello World!'

2_example.py



Optimization Procedure



```
35 # Optimization
36 optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.5)
37 gradients = optimizer.compute_gradients(cross_entropy)
38 train_step = optimizer.apply_gradients(gradients)
```

'Hello World!'

Running computational graph




```
49 #####
50 ##### RUNNING THE COMPUTATIONAL GRAPH #####
51 #####
52 with tf.Session() as sess:
53     writer = tf.summary.FileWriter('./2e', sess.graph) # TensorBoard writer
54     # We need to initialize the variables
55     tf.global_variables_initializer().run()
56
57     # Train the model
58     for i in range(max_iterations):
59         batch_xs, batch_ys = mnist.train.next_batch(batch_size)
60         _, loss, summ = sess.run([train_step, cross_entropy, merged],
61                                 feed_dict={x_in: batch_xs, y_in: batch_ys})
62         writer.add_summary(summ, i)
63
```

'Hello World!'


2_example.py



accuracy



loss





ACCURACY: 89,73%

'Hello World!'

3_example.py



Let's go deeper...



'Hello World!'

3_example.py



```
15 # Declare parameters #
16 weights = {
17     'w1': tf.Variable(tf.random_normal([784, 256])),
18     'w2': tf.Variable(tf.random_normal([256, 128])),
19     'w3': tf.Variable(tf.random_normal([128, 10]))
20 }
21 biases = {
22     'b1': tf.Variable(tf.random_normal([256])),
23     'b2': tf.Variable(tf.random_normal([128])),
24     'b3': tf.Variable(tf.random_normal([10]))
25 }
26
```


```
31 # Inference. Compute everything until the predictions
32 y = tf.nn.sigmoid(tf.matmul(x_in, weights['w1']) + biases['b1']) # INPUT LAYER
33 y = tf.nn.sigmoid(tf.matmul(y, weights['w2']) + biases['b2']) # HIDDEN LAYER 1
34 y = tf.matmul(y, weights['w3']) + biases['b3'] # HIDDEN LAYER 2
35 y_pred = tf.nn.softmax(y) # PREDICTION
```

'Hello World!'


3_example.py



accuracy



loss




ACCURACY: 92.90%

'Hello World!'

4_example.py



We can build autoencoders too!



'Hello World!'

4_example.py



```
15 # Declare parameters #
16 weights = {
17     'encoder_w1': tf.Variable(tf.random_normal([784, 256])),
18     'encoder_w2': tf.Variable(tf.random_normal([256, 64])),
19     'decoder_w1': tf.Variable(tf.random_normal([64, 256])),
20     'decoder_w2': tf.Variable(tf.random_normal([256, 784]))
21 }
22 biases = {
23     'encoder_b1': tf.Variable(tf.random_normal([256])),
24     'encoder_b2': tf.Variable(tf.random_normal([64])),
25     'decoder_b1': tf.Variable(tf.random_normal([256])),
26     'decoder_b2': tf.Variable(tf.random_normal([784]))
27 }
28
```

'Hello World!'

4_example.py



```
29 # Building the encoder
30 def encoder(x):
31     # Encoder Hidden layer with sigmoid activation #1
32     layer_1 = tf.nn.sigmoid(tf.add(tf.matmul(x, weights['encoder_w1']),
33                                     biases['encoder_b1']))
34     # Encoder Hidden layer with sigmoid activation #2
35     layer_2 = tf.nn.sigmoid(tf.add(tf.matmul(layer_1, weights['encoder_w2']),
36                                     biases['encoder_b2']))
37     return layer_2
38
39
40 # Building the decoder
41 def decoder(x):
42     # Decoder Hidden layer with sigmoid activation #1
43     layer_1 = tf.nn.sigmoid(tf.add(tf.matmul(x, weights['decoder_w1']),
44                                     biases['decoder_b1']))
45     # Decoder Hidden layer with sigmoid activation #2
46     layer_2 = tf.nn.sigmoid(tf.add(tf.matmul(layer_1, weights['decoder_w2']),
47                                     biases['decoder_b2']))
48     return layer_2
49
```

'Hello World!'



4_example.py



```
50 # Inputs
51 x_in = tf.placeholder(tf.float32, [None, 784])
52
53 # Inference. Compute everything until the predictions
54 encoder_op = encoder(x_in)
55 decoder_op = decoder(encoder_op)
56 y_pred     = decoder_op
57
58 # Cost Function. Compare predictions with true labels
59 # In this case with use the L2 distance
60 loss = tf.reduce_mean(tf.pow(x_in - y_pred, 2))
61
62 # Optimization
63 optimizer = tf.train.RMSPropOptimizer(learning_rate=0.01)
64 gradients = optimizer.compute_gradients(loss)
65 train_step = optimizer.apply_gradients(gradients)
66
```


'Hello World!'

4_example.py



Done!

Caffe
theano
PYTORCH



DL4J
DEEPLARNING4J